

LSTM

✓ Here's what you're covering and why it works well:

Part	<u>What You're Teaching</u>	<u>Why It's Important</u>
1	Brief on ANN	Sets the base — students know how traditional networks work
2	Problem with ANN for sequence data	Shows the need for RNNs (no memory, fixed input size)
3	Introduction to RNN	Introduces recurrence, hidden state, time step-based processing
4	Problems with RNN (e.g., vanishing gradients)	Prepares the ground for why LSTM was invented
5	Introduction to LSTM	Logical next step — introduce as a solution
6	Basic LSTM Architecture	Gives structural understanding — memory cell + gates
7	Overview of Forget Gate, Input Gate, Output Gate	Core components — gives interpretability and depth

Introduction to LSTM (Long Short-Term Memory)

What is LSTM?

LSTM is a special type of **Recurrent Neural Network (RNN)** designed to **remember information for long periods**.

- Introduced by **Hochreiter & Schmidhuber (1997)**
 - It solves key issues in traditional RNNs such as:
 - **Vanishing gradient**
 - **Difficulty in learning long-term dependencies**
-

Why was LSTM needed?

Let's first understand the problems with RNNs.

Problems in RNN

1. Short-Term Memory

- RNNs can only remember **recent** inputs.
 - Struggle with remembering **long-range dependencies** (e.g., "The boy who wore a red hat... was playing.")
-

2. Vanishing Gradient Problem

- During backpropagation through time (BPTT), gradients are **repeatedly multiplied by weights < 1** .
 - This causes the gradients to shrink exponentially.
 - Eventually, **weights stop updating** → the model forgets earlier information.
-

3. Exploding Gradient Problem (less common)

- If weights > 1 , gradients grow too large → unstable training.
-

4. Difficulty in Capturing Long-Term Context

- Language and time series often require remembering context far back (e.g., meaning of a sentence).
 - RNNs just can't **bridge that gap** effectively.
-



How LSTM Solves These Problems

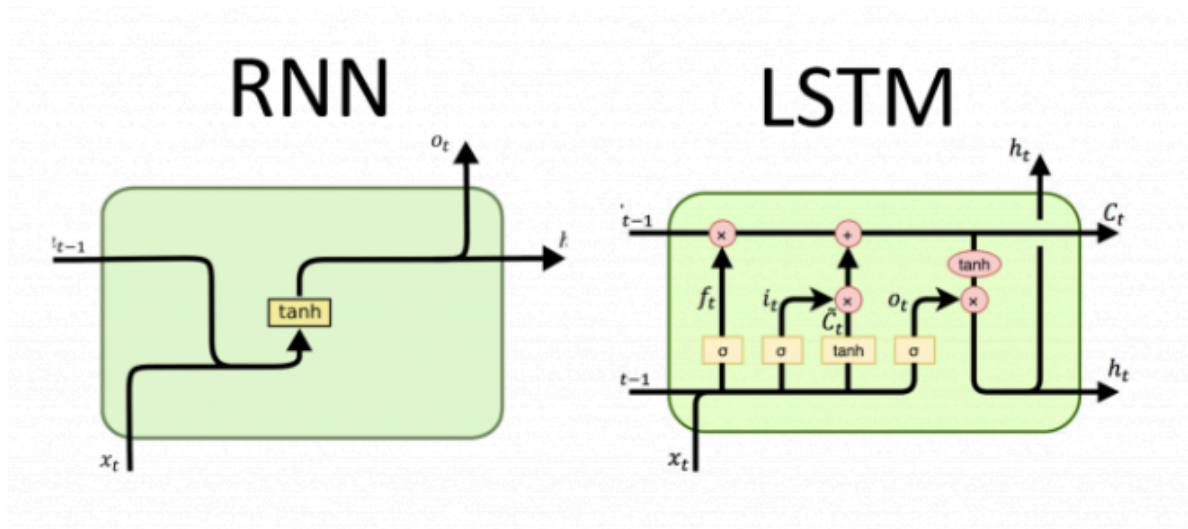
Core Idea:

Instead of having just one hidden state like RNN, LSTM introduces a **Cell State (Ct)** — a "**conveyor belt**" that carries information through time steps **with minimal modification**.

This lets LSTM **retain long-term memory**.



LSTM Architecture Components



LSTM uses **three gates** to regulate the flow of information:

1. Forget Gate (f_t)

- **Decides what to forget** from the previous cell state.
 - Looks at h_{t-1} and x_t , and outputs a number between 0 and 1.
 - 0 = forget completely, 1 = keep everything.
-

2. Input Gate (i_t)

- **Decides what new information** to store in the cell state.
 - Two parts:
 - A sigmoid layer: which values to update
 - A tanh layer: creating candidate values
-

3. Cell State Update

This line is the **heart of LSTM**.
It controls **what to forget + what to add** into memory.

4. Output Gate (ot)

- **Decides what part of the memory to output**
-

Summary Table:

Problem in RNN	How LSTM Solves It
Can't remember long-term info	Cell state carries long-term memory
Vanishing gradient	Additive updates to cell state, avoids exponential shrinkage
Forgetting earlier inputs	Forget gate learns to retain or discard information
Inflexible memory update	Input + Output gates allow precise memory control

Intuition Recap

- **Forget gate:** What to throw away?
- **Input gate:** What to store?
- **Output gate:** What to pass to the next time step?

Think of it like a **smart memory box** with gates deciding what goes in, what stays, and what comes out.

